
Predicting Global Sales of Video Games

Kevin Weng, Yi Lyu, Chen Li, Omar Hafez
MATH 156 Final Project: Group 5
University of California, Los Angeles (UCLA)
[Github Project](#)

Abstract

Our project aims to model global video game sales by applying a neural network, Random Forest, and k-Nearest Neighbors model to preprocessed data to successfully predict global sales and determine which model works best. We built the project on Python using sklearn, Keras, and Tensorflow for the models and SQL for data processing. We will evaluate the results by root mean squared error (RMSE). We hope to see whether or not global sales can be effectively modeled and also give an analysis on model strengths and weaknesses for this task.

1 Introduction

The expansion of the video game industry and eSports in the past decade has fueled gamers and game studios all over the world. With the backdrop of COVID-19, worldwide quarantines, and work-from-home structures, video games have garnered another boost in popularity and success in 2020. With companies like Nintendo, Activision Blizzard, and many independent developers making an impact on modern culture, the landscape of video games and entertainment software has drastically changed. Creating a video game is an intensive project that requires a diverse array of resources and specialists that would be very costly for the studio if a release goes awry. For video game producers, investors, and consultants, the ability to project global sales is very useful when it comes to considering possible translations, global releases, and general marketing investment in other parts of the world.

In this work, we apply artificial neural network, Random Forest, and k-Nearest Neighbors to model global video game sales in Section [2], select the optimal prediction model in Section [5.3], and discuss the predictability of video game sales based on our data in Section [6]. Before describing the sales prediction models, we first provide an overview of related works in Section [2] and the preprocessing steps taken on the dataset in Section [3.2].

2 Related Works

There are other works related to our project in the video game industry, and countless more in the generalized sales and marketing prediction field. One of the projects involves internet search volume as a feature to predict global sale and is becoming increasingly relevant as social media dominates the information space in the majority of the video game industry's target audience^[5]. This data is likely heavily correlated to global sales as the consumer sentiment is captured even prior to release. Another paper that used neural networks predicted weekly game sales on PCA preprocessed data^[3]. The weekly timeframe is different from our cumulative global sales number and may be impacted seasonally. Additionally, one other paper used sexualized cover art content as a feature to predict sales for video games^[4]. In this case, another specific feature was analyzed that we were not able to consider for our project. As pertaining to sales, features regarding behavioral economics, consumer psychology, and marketing can all be possible candidates to more successfully model video game sales.

3 Dataset and Features

The training data set is **Video Game Sales with Ratings** from Kaggle. The dataset consists of 11,563 video game titles detailing release year, publisher, platform, genre, regional sales, global sales, critic and user scores, critic and user counts, and ESRB rating. Not all features are present for every title. The critic and user scores were obtained from Metacritic, a popular video game review site.

| | Platform | Year_of_Release | Genre | Publisher | Global_Sales | Critic_Score | Critic_Count | User_Score | User_Count | Developer | Rating |
|---|----------|-----------------|--------------|-----------|--------------|--------------|--------------|------------|------------|-----------|--------|
| 0 | Wii | 2006.0 | Sports | Nintendo | 82.53 | 76.0 | 51.0 | 8 | 322.0 | Nintendo | E |
| 1 | NES | 1985.0 | Platform | Nintendo | 40.24 | NaN | NaN | NaN | NaN | NaN | NaN |
| 2 | Wii | 2008.0 | Racing | Nintendo | 35.52 | 82.0 | 73.0 | 8.3 | 709.0 | Nintendo | E |
| 3 | Wii | 2009.0 | Sports | Nintendo | 32.77 | 80.0 | 73.0 | 8 | 192.0 | Nintendo | E |
| 4 | GB | 1996.0 | Role-Playing | Nintendo | 31.37 | NaN | NaN | NaN | NaN | NaN | NaN |

Figure 1: First 5 entries of dataset

3.1 Features

- Name** - The name of the video game.
- Platform** - The console on which the game runs on. (Wii, PS4, PC, etc.)
- Year of Release** - The year the game was released.
- Genre** - Category of the game. (Shooter, Racing, Puzzle, etc.)
- Publisher** - Publisher of the game.
- NASales, EUSales, JPSales, OtherSales** - regional sales of video games in millions of units.
- Global Sales** - Total sales in the world in millions of units on a particular platform.
- Critic score** - Score by Metacritic's critics.
- Critic count** - Number of critics who contributed to Critic_score.
- User score** - Score by Metacritic's subscribers.
- User count** - Number of users who contributed to User_score.
- Developer** - Party who created the game.
- Rating** - The Entertainment Software Rating Board (ESRB) rating.

3.2 Preprocessing

First, we removed the games missing **Platform, Genre, Publisher, and Year of Release** data in that these variables could not be imputed effectively. We then imputed with median **Critic score, User score, Critic count, and User count** because many are missing values. For the remaining categorical data in **Genre and Publisher**, we used one-hot-encoding to make the data suitable for regression and dropped one column of each to decorrelate the columns. Our preprocessing also removed the local sales of games because they were often obtained after global sales were calculated and thereby would not be useful for showing a correlation with prior regional video game sales and a global launch. Because of the uncommon popularity of certain games, such as Grand Auto Theft V, we consider the top 10% and bottom 10% as outliers and remove them accordingly. Also, games published before 2004 are not considered due to a smaller size of the gaming industry at that time; the year 2004 is chosen for the release of War of Warcraft (WoW). In addition, we group video games by their **names** so that the global sales data do not concern the platforms these games are/were published on.




Figure 2: Correlation Matrix of Features (Partial)

After preprocessing the data, we split the remaining into 80% training data and 20% test data. For k-Nearest Neighbors, we use cross validation to determine the optimal number of nearest neighbors. Two approaches have been attempted: we split the data into three categories: training, test, and cross validation, which is then used to determine the optimal k number of neighbors that minimizes RMSE; we apply k -fold cross validation to the dataset, which will be discussed in detail in Section 4.1. The correlation matrix in Figure 2 is used to determine whether or not any features should be dropped. After feature evaluation, we find that **critic score, user score, genre, and publisher** potentially have a great impact on video games sales prediction and are consequently chosen to be the features we use for training.

4 Methods

4.1 Models

Neural Network - The Neural Network, as studied in class, is a model comprised of hidden layers of nodes known as neurons that takes in an input and, by feeding it through the hidden layers, produces a result in the output layer. In the Neural Network model we use, we apply the state-of-the-art optimizer NAdam^[7] to our model and use the rectified linear unit (ReLU) as the activation function in that functions like tanh and sigmoid are hard to train given the limited size of our dataset, and ReLU, by definition, does not permit negative values, which coincides with our purpose of predicting global sales.

Random Forest - The Random Forest model is an ensemble method that trains multiple decision trees and outputs a class by majority vote^[1]. For regression tasks, instead of mode, the mean prediction of the individual trees is returned. For reference, a decision tree is a popular machine learning algorithm that uses many input variables to traverse down a tree, and

returns a prediction from a leaf. The benefit of using multiple trees is a reduced variance as a single tree can easily overfit the data. Random Forest differs from simply bagging multiple decision trees by selecting a random subset of the features for each tree so that if some features are stronger predictors than others, such trees would be correlated known as the 'Random Subspace method'^[6].

k-Nearest Neighbors - The k-Nearest Neighbors algorithm takes an element and looks for its closest neighbors to take a majority vote in the classification case, and the mean or median value of the nearest neighbors for regression^[2]. Mean is chosen in our case. The basic structure of the algorithm is described as follows: for all possible k for our model,

1. We divide the training dataset into p equal parts, where p is fixed.
2. We randomly choose one part for cross validation and the remaining $p - 1$ parts for training, which we will repeat for p times so that each part is used once as cross validation set, yielding us p errors. We then compute the average error $\epsilon(k)$ of this model given k over the p parts.
3. We find the k that minimizes the average error $\epsilon(k)$ and return the model.

For our experiment, we weighed the points by the inverse of their distance, making closer neighbors have greater influence: $\hat{y} = \sum_{i=1}^k \frac{d(x, x_i) y_i}{\sum_{i=1}^k d(x, x_i)}$

5 Experiments/Results/Discussion

5.1 Methodology

After training the three models, we will evaluate the results with the root mean squared error (RMSE) metric to evaluate efficacy and accuracy.

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

where y_i is the global sales corresponding to the test input x_i , and \hat{y}_i is the prediction of our models. RMSE takes into account negative values and is a commonly used metric in determining the performance of regression models.

5.2 Results

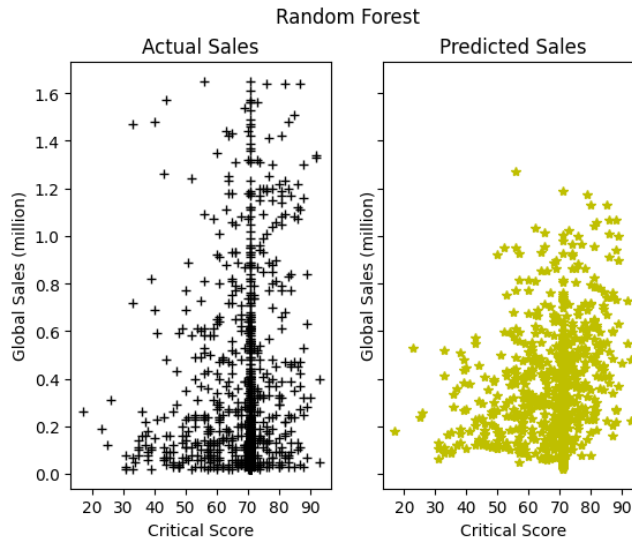


Figure 3: Random Forest Predictions




Figure 4: k-Nearest Neighbors Predictions

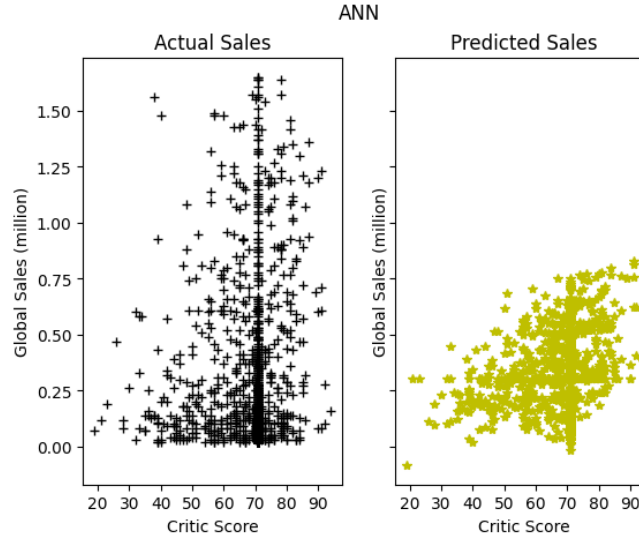


Figure 5: Artificial Neural Network Predictions

As we can see in Figure 3, Figure 4, and Figure 5¹, the predicted sales roughly capture the general distribution of the global sales, but the one corresponding to k nearest neighbors regression is more localized than are the other two, which can be ascribed to the fact that the k nearest neighbors regression only looks at the neighbors of the input for prediction, while the other two consider the data globally. Although Artificial Neural Network appears better than KNN in predicting global sales, it returns a negative global sales, which is undesirable. Through computation, we can have

| | |
|-----------------------|--------|
| RMSE _{RF} : | 0.3216 |
| RMSE _{kNN} : | 0.3315 |
| RMSE _{ANN} : | 0.3562 |

¹The labels of x axis in Figure 3, Figure 4, and Figure 5 above are misleading because features other than critic score have been used in our models; we consider scatter plots are better in terms of visualization.

5.3 Discussion

From our RMSE values, we see that the strongest performer was Random Forest and the weakest was the ANN trained with 10 epochs. Increasing the number of epochs to 200 improved the RMSE to 0.3319, still the worst performing model. Naturally, we will discuss the strengths and weaknesses of our models. For Random Forest, an ensemble of regression trees, the individual trees are easy to understand, specifically for mirroring human behavior purposes which is relevant to our task of suggesting that features impact decisions to purchase video games, but because of the nature of Random Forest, modelers have few controls over it. In contrast with Random Forest, Neural Network is incomprehensible, making it less attractive for practical uses of sales prediction. However, as mentioned previously, there have been neural networks applied to PCA processed data for sales prediction. Regarding the k -Nearest Neighbors algorithm, as mentioned in Section 5.2 above, it falls in between the other two models in terms of RMSE, which implies that the model performs relatively well on this dataset, but the predicted sales for k -Nearest Neighbors regression are localized, as shown in Figure 4, which leads to a low variance. In addition, that k -Nearest Neighbors regression has to store all the data in the training set renders it inefficient when the size of the data set increases.

6 Conclusion and Future Work

With an RMSE of around 0.33, our objective of modeling global sales of video games is decently accurate and can provide some insight for future video game releases, especially considering that the range of global sales value is 0 to 60. Also, the models have a low chance of overfitting because ensemble methods and cross validation, which are used in our models, are innately preventative measures against overfitting. Out of the regressors that we explored, we conclude that the Random Forest model produced the best prediction of global video game sales by measure of RMSE, albeit by a small margin. To continue our exploration in the realm of video games, we see that some of the most successful and popular games today are free to play, offering paid in-game content that users can elect to pay for or not. Our project does not strongly consider this model of games and exploration of different payment structures may be interesting to consider in the future. Plus, because we exclude those extremely popular and unpopular as outliers, our models can only predict the sales of normal video games. Furthermore, as seen in related work, there are many esoteric features that were not taken into account. As with any consumer product, sentiment is a major factor in how a product is received and could also be a viable direction of exploration. Separate analyses of games and marketing strategies can be helpful to building a more holistic and complete model of global sales.

References

- [1] Anava, Oren, and Kfir Levy. "k*-nearest neighbors: From global to local." Advances in neural information processing systems. 2016.
- [2] Breiman, Leo. "Random forests." Machine learning 45.1 (2001): 5-32.
- [3] Marcoux, Julie, and Sid-Ahmed Selouani. "A hybrid subspace-connectionist data mining approach for sales forecasting in the video game industry." 2009 WRI World Congress on Computer Science and Information Engineering. Vol. 5. IEEE, 2009.
- [4] Near, Christopher E. "Selling gender: Associations of box art representation of female characters with sales for teen-and mature-rated video games." Sex roles 68.3 (2013): 252-269.
- [5] Ruohonen, Jukka, and Sami Hyrynsalmi. "Evaluating the use of internet search volumes for time series modeling of sales in the video game industry." Electronic Markets 27.4 (2017): 351-370.
- [6] Wikipedia contributors. "Random subspace method." Wikipedia, The Free Encyclopedia. Wikipedia, The Free Encyclopedia, 3 Nov. 2019. Web. 27 Jul. 2020.
- [7] Dozat, Timothy. "Incorporating Nesterov Momentum into Adam." ICLR 2016 workshop. 2016.

```

167 """
168 =====
169 == Filename: helper.py
170 == Author: Yi Lyu
171 == Status: Complete
172 =====
173 """
174
175 import numpy as np
176 import pandas as pd
177 import pickle
178 import sqlite3
179 import re
180 import os
181 from sklearn.impute import SimpleImputer as Imputer
182 from sklearn.impute import KNNImputer
183 from sklearn.preprocessing import LabelEncoder
184
185 __all__ = ['Videogames']
186
187 def get_dir(path):
188     return os.path.join(getWorkDir(), path)
189
190 def getWorkDir():
191     pathlist = os.path.abspath(os.curdir).split('/')
192     path = '/'
193     for p in pathlist:
194         path = os.path.join(path, p)
195         if p == 'video-game-sales-predictor' or p == 'video-game-sales-
196 predictor-master':
197             break
198     return path
199
200 class Videogames(object):
201     def __init__(self, database_dir, data_dir='data/', storage='data'):
202         self.database_dir = database_dir
203         self.table = ''
204         self.data_dir = data_dir
205         self.storage = '{0}.pickle'.format(storage)
206
207         self._has_data = False
208         self._headers = []
209         self._dtypes = []
210         self._connection = None
211         try:
212             with open(get_dir(data_dir + self.storage), "rb") as f:
213                 self.table, self._headers, self._dtypes, self._has_data =
214 pickle.load(f)
215         except:
216             pass
217
218     @property
219     def table_name(self):
220         return self.table
221
222     @property
223     def status(self):
224         return self.get_status()
225
226     @property
227     def headers(self):
228         return self._headers
229
230     @property
231     def dtypes(self):

```



```

232         return self._dtypes
233
234     def get_status(self):
235         return self._has_data
236
237     def get_headers(self):
238         return self._headers
239
240     def get_dtypes(self):
241         return self._dtypes
242
243     def read_data_in(self, filepath, table, write_headers=False):
244         conn = sqlite3.connect(database=self.database_dir)
245         cur = conn.cursor()
246         data = pd.read_csv(filepath, delimiter=",", encoding="
247 unicode_escape")
248
249         self.table = table
250         headers = self._get_headers(data)
251         dtypes = self._get_dtypes(data)
252         self._create_table(headers, dtypes, cur)
253
254         if write_headers:
255             with open(get_dir(self.data_dir + 'headers.csv'), "w+") as f:
256                 f.write(", \n".join(headers))
257
258         if not self._has_data:
259             data = self._remove_missing(data)
260             with open(get_dir(self.data_dir + self.storage), "wb+") as f:
261                 self._insert_data(data, headers, dtypes, cur)
262                 self._has_data = True
263                 pickle.dump((self.table, headers, dtypes, True), f,
264 pickle.HIGHEST_PROTOCOL)
265
266         del data
267         conn.commit()
268         conn.close()
269
270     def _remove_missing(self, data):
271         data = data.replace(r'tbd', np.nan, regex=True)
272         data['User_Score'] = data['User_Score'].astype(np.float64)
273         condition = (data['Platform'].notnull() & data['Genre'].notnull()
274 & data['Publisher'].notnull() & data['Year_of_Release'].notnull())
275         data = self._imputation(data[condition])
276         return data
277
278     def _imputation(self, data):
279         imp = Imputer(strategy='median')
280         attributes = ['Critic_Score', 'User_Score', 'Critic_Count', '
281 User_Count']
282         for item in attributes:
283             data[item] = imp.fit_transform(data[[item]]).ravel()
284         return data
285
286     def get_col(self, *header):
287         if not self._connection:
288             self._connection = sqlite3.connect(self.database_dir)
289             cur = self._connection.cursor()
290
291             command = "SELECT {0} FROM {1};".format(self._list2str(header),
292 self.table)
293             return self._col2list(cur.execute(command).fetchall())
294
295     def execute(self, command):
296         if not self._connection:

```



```

297         self._connection = sqlite3.connect(self.database_dir)
298         cur = self._connection
299
300         if bool(re.match("^[\t\n]*SELECT", command, re.I)):
301             return list(self._col2list(cur.execute(command).fetchall()))
302         else:
303             print("ILLEGAL COMMAND")
304
305
306     ## Helper Functions ##
307     def _get_headers(self, data):
308         """Return the headers of the data
309
310         Args:
311             data DataFrame: the data we read from csv.
312
313         Returns:
314             list: the headers of the data
315         """
316         if not self._headers:
317             self._headers = list(map(lambda col: col.lower(), data.
318 columns))
319         return self._headers
320
321     def _get_dtypes(self, data):
322         if not self._dtypes:
323             self._dtypes = [self._process_dtype(data[col][0]) for col in
324 data.columns]
325         return self._dtypes
326
327     def _create_table(self, headers, dtypes, cur):
328         """Execute the following SQL command
329
330         CREATE TABLE IF NOT EXISTS {table} (
331             name VARCHAR(80),
332             ...
333         );
334
335         Args:
336             headers (list): the list of columns where each header is
337 lowercase.
338             dtypes (list): the list of types where each type is either
339 NUMBER or VARCHAR(80) based on this data set.
340             cur (sqlite3.connection.cursor): a connection cursor of
341 sqlite3 database
342         """
343         command = "CREATE TABLE IF NOT EXISTS {0} {}".format(self.table)
344         template = "{0} {1}"
345
346         n = len(headers)
347
348         ## Convert the data to suitable form for _list2str function
349         data = [template.format(headers[i], dtypes[i]) for i in range(n)]
350
351         command += self._list2str(data)
352         command += ";"
353         cur.execute(command)
354
355     def _insert_data(self, data, headers, dtypes, cur):
356         command_template = "INSERT INTO {0} ({1}) VALUES ({2});"
357         for i, itr in data.iterrows():
358             res = list(map(self._str_classifier, list(itr)))
359             command = command_template.format(self.table, ", ".join(
360 headers),

```

```

361                                     self._list2str(res,
362 classify=self._row_classifier(res, dtypes))
363                                     cur.execute(command)
364
365 def _list2str(self, data, delimiter=",", classify=lambda x: x):
366     """Convert the list to a string
367
368     I have not found such a function in Python and therefore
369     wrote one.
370
371     Args:
372         data (list): the row of the table
373         delimiter (str, optional): the delimiter.
374         classify (function, optional): a function that classifies the
375         data in the row.
376
377     Returns:
378         str: a string representing the data converted to a string.
379     """
380     res = ""
381     for i in range(len(data)):
382         res += classify(data[i])
383         if i != len(data) - 1:
384             res += delimiter + " "
385     return res
386
387 def _row_classifier(self, data, dtypes):
388     """ classify the data in a row in the table
389     def classifier(x):
390         i = data.index(x)
391         if dtypes[i] == "NUMBER":
392             if x == "NULL" or x == 'tbd':
393                 return "-1"
394             else:
395                 return str(x)
396         else:
397             return "{}\{}".format(x)
398     return classifier
399
400 def _str_classifier(self, x):
401     """ classify the data so that it does not contain nan
402     if type(x) == float and np.isnan(x):
403         return -1
404     return x
405
406 def _process_dtype(self, var):
407     dtype = type(var)
408     if dtype == str and var.isnumeric():
409         return "NUMBER"
410     type_converter = {type(','): "VARCHAR(80)", np.float64: "NUMBER"
411 , np.int64: "NUMBER"}
412     return type_converter[dtype]
413
414 def _col2list(self, col):
415     n = len(col[0])
416     return list(map(lambda x: list(x)[:n], col))
417
418 """
419 ==
420 == Filename: main.py
421 == Author: Yi Lyu
422 == Status: Complete
423 ==
424 """

```

```

425 import numpy as np
426 import matplotlib.pyplot as plt
427 import pandas as pd
428 import seaborn as sns
429 import pickle
430 import os
431 from sklearn.decomposition import PCA
432 from keras.models import load_model
433 from sklearn.model_selection import train_test_split
434
435 from helper import Videogames, getWorkDir, get_dir
436 from models import *
437 from plotting import *
438
439 def read_data():
440     videogames = Videogames(get_dir("data/math156.db"))
441     videogames.read_data_in(get_dir("data/videogames.csv"), "VIDEOGAMES",
442                             True)
443     res = np.array(videogames.execute('''
444         SELECT name, g_total, cscore, uscore, genre, publisher FROM (
445             SELECT name AS name,
446                   SUM(global_sales) AS g_total,
447                   critic_score AS cscore,
448                   user_score AS uscore,
449                   genre AS genre,
450                   publisher AS publisher
451             FROM VIDEOGAMES
452             WHERE year_of_release >= 2004 and uscore != 0 and cscore != 0
453             GROUP BY name) AS VideogameSummary
454         WHERE g_total != 0
455         ORDER BY g_total DESC;
456     '''))
457     return res
458
459 if __name__ == "__main__":
460     ## the critic scores and user scores
461     columns = ['name', 'gtotal', 'cscore', 'uscore', 'genre', 'publisher'
462               ]
463     res = pd.DataFrame(read_data(), columns=columns)
464
465     n = len(res)
466     factor = 0.1
467     quantile1 = round(n * factor)
468     quantile2 = n - round(n * factor)
469     res = res.loc[quantile1:quantile2 + 1, :]
470
471     ## Transform data into appropriate form for regression
472     scores = res[['cscore', 'uscore']]
473     genre = pd.get_dummies(res['genre'], drop_first=True)
474     publisher = pd.get_dummies(res['publisher'], drop_first=True)
475
476     X = pd.concat((scores, genre, publisher), axis=1).astype('float64')
477     Y = res['gtotal'].astype('float64')
478
479
480     X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=
481     .20, train_size=.80, random_state = 40)
482
483     try:
484         with open(get_dir('data/models.pickle'), 'rb') as f:
485             rfreg, knnreg = pickle.load(f)
486             annreg = load_model(get_dir('data/ann'))
487     except:
488         annreg = ann(X_train, Y_train.ravel()) ## ANN

```

```

489         rfregr = random_forest(X_train, Y_train.ravel())    ## Random
490     Forest
491         knnregr = knn(X_train, Y_train.ravel())             ## KNN
492         with open(get_dir('data/models.pickle'), 'wb+') as f:
493             pickle.dump((rfregr, knnregr), f, pickle.HIGHEST_PROTOCOL)
494             annregr.save(get_dir('data/ann'))
495
496     print("The mean is", np.mean(Y))
497     ## RMSE
498     rmse_template='RMSE\t{name:25}{value:18}'
499     print(rmse_template.format(name='random forest', value=rmse(X_test,
500     Y_test, rfregr)))
501     print(rmse_template.format(name='Knn', value=rmse(X_test, Y_test,
502     knnregr)))
503     print(rmse_template.format(name='ANN', value=rmse(X_test, Y_test,
504     annregr)))
505
506     plot_predictions(X_test, Y_test, rfregr, knnregr, annregr)
507
508     """
509
510     print('=====')
511
512     ## R2
513     r2_template = 'R^2\t{name:25}{value:18}'
514     print(r2_template.format(name='random forest', value=rfregr.score(
515     X_test, Y_test)))
516     print(r2_template.format(name='Knn', value=knnregr.score(X_test,
517     Y_test)))
518     print(r2_template.format(name='Aan', value=annregr.score(X_test,
519     Y_test)))
520
521     """
522
523     """
524     =====
525     ==  Filename: models.py
526     ==  Author: Yi Lyu
527     ==  Status: Complete
528     =====
529     """
530
531     import numpy as np
532     import matplotlib.pyplot as plt
533     from sklearn.linear_model import Ridge
534     from sklearn.linear_model import GammaRegressor
535     from sklearn.preprocessing import PolynomialFeatures, StandardScaler
536     from sklearn.pipeline import make_pipeline
537     from sklearn.ensemble import RandomForestRegressor
538     from sklearn.neighbors import KNeighborsRegressor
539     from sklearn.model_selection import train_test_split, cross_val_score
540     from sklearn.metrics import mean_squared_error
541     from keras.models import Sequential
542     from keras.layers import Dense
543
544     ## just in case someone wants to implement them instead of using sklearn
545
546     def rmse(X_test, Y_test, model):
547         Y_pred = model.predict(X_test)
548         return mean_squared_error(Y_test, Y_pred, squared=False)
549
550     def plot_knn(ns, rmses):
551         plt.plot(ns, rmses, 'r*')
552         plt.xlabel('# of neighbors')

```

```

553 plt.ylabel('RMSE')
554
555 plt.savefig('graphs/knn_choice_n.png', bbox_inches='tight')
556 plt.clf()
557
558 def knn(xs, ys, n=10):
559     #X_train, X_test, Y_train, Y_test = train_test_split(xs, ys,
560     test_size= .1, random_state = 40)
561     num_cols = len(xs.columns)
562     i = 5
563
564     best_index = 4
565     best_score = 10000
566     nums = [i for i in range(5, int(np.sqrt(num_cols)) + 10)]
567     cvs = []
568
569     for num in nums:
570         model = KNeighborsRegressor(n_neighbors=num, algorithm='kd_tree',
571         weights='distance')
572         temp = cross_val_score(model, xs, ys, cv=5).mean()
573         temp = np.sqrt(1 - temp)
574         if temp < best_score:
575             best_score = temp
576             best_index = num
577     print(best_index)
578     return KNeighborsRegressor(n_neighbors=best_index, algorithm='kd_tree',
579     weights='distance').fit(xs, ys)
580     """
581
582     best_model = KNeighborsRegressor(n_neighbors=i, algorithm='kd_tree',
583     weights='distance').fit(X_train, Y_train)
584     best_rmse = rmse(X_test, Y_test, best_model)
585
586     ### Cross Validation
587     ns = [n]
588     rmses = [best_rmse]
589     cvs = []
590     ### You can change 5 to * 2 or * 3 here for a better result, but
591     slower.
592     for n in range(i, int(np.sqrt(num_cols)) + 5):
593         model = KNeighborsRegressor(n_neighbors=n, algorithm='kd_tree',
594         weights='distance').fit(X_train, Y_train)
595         temp = rmse(X_test, Y_test, model)
596         ns.append(n)
597         rmses.append(temp)
598         if temp < best_rmse:
599             best_model = model
600             best_rmse = temp
601     plot_knn(ns, rmses)
602
603     """
604
605     return best_model
606
607 def ann(xs, ys):
608     n = len(xs.columns)
609     ANN = Sequential()
610     ANN.add(Dense(units = 6, activation = "relu", input_dim = n))
611     ANN.add(Dense(units = 4, activation = "relu"))
612     ANN.add(Dense(units = 1))
613
614     ANN.compile(optimizer = "adam", loss = "mean_squared_error")
615     ANN.fit(xs, ys, batch_size = 2, epochs = 200)
616     return ANN
617

```

```

618 def gamma_model(xs, ys):
619     model = GammaRegressor().fit(xs, ys)
620     return model
621
622 def linear_model(xs, ys, m):
623     model = make_pipeline(PolynomialFeatures(m), Ridge(normalize=True)).
624     fit(xs, ys)
625     return model
626
627 def random_forest(xs, ys):
628     model = RandomForestRegressor(criterion='mse').fit(xs, ys)
629     return model

630 """
631 =====
632 == Filename: plotting.py
633 == Author: Yi Lyu
634 == Status: Complete
635 =====
636 """
637
638 import numpy as np
639 import matplotlib.pyplot as plt
640 import pandas as pd
641 import seaborn as sns
642 import os
643
644 from helper import getWorkDir, get_dir
645
646 def predict(X_test, Y_test, model):
647     """Predict the sales based on the dataset
648
649     Args:
650         X_test (DataFrame): Data
651         Y_test (Series): Actual Sales
652         model (object): Model we are using
653
654     Returns:
655         DataFrame: predicted scales
656     """
657     return pd.DataFrame(model.predict(X_test))
658
659 def plot_helper(xs, data_ys, predict_ys, model_name='Unknown'):
660     """Plot the predicted sales
661
662     Args:
663         xs (Series): the x values
664         data_ys (Series): the actual sales
665         predict_ys (Series): the predicted sales
666         model_name (str, optional): the name of the model. Defaults to '
667     Unknown'.
668     """
669     xs = xs.astype(np.float64)
670     fig, (ax1, ax2) = plt.subplots(1, 2, sharex=True, sharey=True)
671     plt.xticks(np.linspace(0, 100, 11))
672
673     fig.suptitle(model_name)
674
675     ax1.plot(xs, data_ys, 'k+', label='data')
676     ax1.set_title('Actual Sales')
677     ax1.set_xlabel='Critic Score', ylabel='Global Sales (million)'
678
679     ax2.plot(xs, predict_ys, 'y*', label='prediction')
680     ax2.set_title('Predicted Sales')
681     ax2.set_xlabel='Critic Score', ylabel='Global Sales (million)'

```

```

682 pic_path = 'graphs/{0}.png'.format(model_name.replace(' ', '_').lower()) 52
683 () 53
684
685 pic_dir = get_dir(pic_path) 54
686
687 plt.savefig(pic_dir, bbox_inches='tight') 55
688 plt.clf() 56
689
690
691 def plot_helper2(data_ys, predicted_ys, model_name='Unknown'): 57
692     fig, (ax1, ax2) = plt.subplots(1, 2, sharex=True, sharey=True) 58
693     fig.suptitle(model_name) 59
694
695     bins = np.arange(0, 6, 0.1) 60
696     sns.distplot(data_ys, bins=bins, hist=True, kde=True, ax=ax1, color='r', 61
697                 axlabel='Sales') 62
698     ax1.set_title('Actual Sales -- Density Plot') 63
699     ax1.set_xlim(0, 2) 64
700     sns.distplot(predicted_ys, bins=bins, hist=True, kde=True, ax=ax2, 65
701                 color='b', axlabel='Sales') 66
702     ax2.set_title('Predicted Sales -- Density Plot') 67
703     ax2.set_xlim(0, 2) 68
704
705     pic_path = 'graphs/{0}_hist.png'.format(model_name.replace(' ', '_'). 69
706     lower()) 70
707     pic_dir = get_dir(pic_path) 71
708
709     plt.savefig(pic_dir, bbox_inches='tight') 72
710     plt.clf() 73
711
712 def plot_predictions(X_test, Y_test, rfregr, knnregr, annregr): 74
713     """Plot the Predicted sales of each model 75
714
715     Args: 76
716         X_test (DataFrame): data 77
717         Y_test (Series): actual sales 78
718         rfregr (RandomForestRegressor): Random Forest Regressor 79
719         knnregr (KNNRegressor): KNN Regressor 80
720         annregr (ANNRegressor): Artificial Neural Network Regressor 81
721     """ 82
722     cscores = X_test['cscore'] 83
723     ## Get predicted sales 84
724     rfres = predict(X_test, Y_test, rfregr) 85
725     knnres = predict(X_test, Y_test, knnregr) 86
726     annres = predict(X_test, Y_test, annregr) 87
727
728     ## Correct the indices in case 88
729     temp = pd.DataFrame(pd.concat([cscores, Y_test], axis=1).to_numpy(), 89
730                         columns=['cscore', 'gtotal'], 90
731                         index=np.arange(0, len(cscores), 1)) 91
732
733     ## Create a pandas DataFrame sorted by Critic Score 92
734     df = pd.concat([temp, rfres, knnres, annres], axis=1) 93
735     df = pd.DataFrame(df.sort_values(by='cscore', ascending=True). 94
736     to_numpy(), 95
737     columns=['cscore', 'gtotal', 'rfres', 'knnres', 'annres'] 96
738     ]) 97
739
740     plot_helper(df['cscore'], df['gtotal'], df['rfres'], 'Random Forest') 98
741     plot_helper(df['cscore'], df['gtotal'], df['knnres'], 'KNN') 99
742     plot_helper(df['cscore'], df['gtotal'], df['annres'], 'ANN') 100
743
744     plot_helper2(df['gtotal'], df['rfres'], 'Random Forest') 101
745     plot_helper2(df['gtotal'], df['knnres'], 'KNN') 102
746     plot_helper2(df['gtotal'], df['annres'], 'ANN') 103

```